

```
In [1]: #获取数据集
import numpy as np
import pandas as pd
from sklearn.datasets import fetch_openml
mnist = fetch_openml('mnist_784', version=1, as_frame=False)
mnist
```

```
'pixel47',
'pixel48',
'pixel49',
'pixel50',
'pixel51',
'pixel52',
'pixel53',
'pixel54',
'pixel55',
'pixel56',
'pixel57',
'pixel58',
'pixel59',
'pixel60',
'pixel61',
'pixel62',
'pixel63',
'pixel64',
'pixel65',
'pixel66'
```


```
In [2]: #查看数组
X, y = mnist["data"], mnist["target"]
X.shape
```

Out[2]: (70000, 784)

```
In [3]: y.shape
```

Out[3]: (70000,)

```
In [4]: %matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
some_digit = X[36000]
some_digit_image = some_digit.reshape(28, 28)
plt.imshow(some_digit_image, cmap = matplotlib.cm.binary,
interpolation="nearest")
plt.axis("off")
plt.show()
```



```
In [5]: y[36000]
```

```
Out[5]: '9'
```

```
In [6]: import matplotlib as mpl

def plot_digit(data):
    image = data.reshape(28, 28)
    plt.imshow(image, cmap = mpl.cm.binary,
                interpolation="nearest")
    plt.axis("off")
def plot_digits(instances, images_per_row=10, **options):
    size = 28
    images_per_row = min(len(instances), images_per_row)
    n_rows = (len(instances) - 1) // images_per_row + 1

    # Append empty images to fill the end of the grid, if needed:
    n_empty = n_rows * images_per_row - len(instances)
    padded_instances = np.concatenate([instances, np.zeros((n_empty, size *
size)), axis=0)

    # Reshape the array so it's organized as a grid containing 28×28 images:
    image_grid = padded_instances.reshape((n_rows, images_per_row, size, size))

    big_image = image_grid.transpose(0, 2, 1, 3).reshape(n_rows * size,
images_per_row * size)
    # Now that we have a big image, we just need to show it:
    plt.imshow(big_image, cmap = mpl.cm.binary, **options)
    plt.axis("off")

plt.figure(figsize=(9,9))
example_images = X[:100]
plot_digits(example_images, images_per_row=10)
plt.savefig("more_digits_plot")
plt.show()
```



```
In [7]: y = y.astype(np.uint8)
X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000], y[60000:]
```

```
In [8]: shuffle_index = np.random.permutation(60000)
X_train, y_train = X_train[shuffle_index], y_train[shuffle_index]
```

```
In [9]: y_train_5 = (y_train == 5) # True for all 5s, False for all other digits.
y_test_5 = (y_test == 5)
```

```
In [10]: from sklearn.linear_model import SGDClassifier
sgd_clf = SGDClassifier(random_state=42)
sgd_clf.fit(X_train, y_train_5)
```

```
Out[10]: SGDClassifier
SGDClassifier(random_state=42)
```

```
In [11]: sgd_clf.predict([some_digit])
```

```
Out[11]: array([False])
```

```
In [12]: from sklearn.model_selection import StratifiedKFold
from sklearn.base import clone
skfolds = StratifiedKFold(n_splits=3, random_state=42, shuffle=True)
for train_index, test_index in skfolds.split(X_train, y_train_5):
    clone_clf = clone(sgd_clf)
    X_train_folds = X_train[train_index]
    y_train_folds = (y_train_5[train_index])
    X_test_fold = X_train[test_index]
    y_test_fold = (y_train_5[test_index])
    clone_clf.fit(X_train_folds, y_train_folds)
    y_pred = clone_clf.predict(X_test_fold)
    n_correct = sum(y_pred == y_test_fold)
    print(n_correct / len(y_pred)) # prints 0.9502, 0.96565 and 0.96495
```

```
0.96775
0.9672
0.9601
```

```
In [13]: from sklearn.model_selection import cross_val_score
cross_val_score(sgd_clf, X_train, y_train_5, cv=3, scoring="accuracy")
```

```
Out[13]: array([0.96055, 0.9585 , 0.96885])
```

```
In [14]: from sklearn.base import BaseEstimator
class Never5Classifier(BaseEstimator):
    def fit(self, X, y=None):
        pass
    def predict(self, X):
        return np.zeros((len(X), 1), dtype=bool)
```

```
In [15]: never_5_clf = Never5Classifier()
cross_val_score(never_5_clf, X_train, y_train_5, cv=3, scoring="accuracy")
```

```
Out[15]: array([0.91275, 0.90845, 0.90775])
```

```
In [16]: from sklearn.model_selection import cross_val_predict
y_train_pred = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3)
```

```
In [17]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_train_5, y_train_pred)
```

```
Out[17]: array([[53952,   627],
               [ 1615,  3806]], dtype=int64)
```

```
In [18]: y_train_perfect_predictions = y_train_5 # pretend we reached perfection
confusion_matrix(y_train_5, y_train_perfect_predictions)
```

```
Out[18]: array([[54579,    0],
               [    0,  5421]], dtype=int64)
```

```
In [19]: from sklearn.metrics import precision_score, recall_score
recall_score(y_train_5, y_train_pred) # == 4344 / (4344 + 1077)
```

Out[19]: 0.7020844862571481

```
In [20]: from sklearn.metrics import f1_score
f1_score(y_train_5, y_train_pred)
```

Out[20]: 0.7724781814491578

```
In [21]: y_scores = sgd_clf.decision_function([some_digit])
y_scores
```

Out[21]: array([-7942.70054026])

```
In [22]: threshold = 0
y_some_digit_pred = (y_scores > threshold)
y_some_digit_pred
```

Out[22]: array([False])


```
In [23]: threshold = 200000
y_some_digit_pred = (y_scores > threshold)
y_some_digit_pred
```

Out[23]: array([False])


```
In [24]: y_scores = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3,
method="decision_function")
```

```
In [25]: from sklearn.metrics import precision_recall_curve
precisions, recalls, thresholds = precision_recall_curve(y_train_5, y_scores)
```

```
In [26]: def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):  
    plt.plot(thresholds, precisions[:-1], "b--", label="Precision")  
    plt.plot(thresholds, recalls[:-1], "g-", label="Recall")  
    plt.xlabel("Threshold")  
    plt.legend(loc="upper left")  
    plt.ylim([0, 1])  
plot_precision_recall_vs_threshold(precisions, recalls, thresholds)  
plt.show()
```



```
In [27]: def plot_precision_vs_recall(precisions, recalls):  
    plt.plot(recalls, precisions, "b-", linewidth=2)  
    plt.xlabel("Recall", fontsize=16)  
    plt.ylabel("Precision", fontsize=16)  
    plt.axis([0, 1, 0, 1])  
    plt.grid(True)  
plt.figure(figsize=(8, 6))  
plot_precision_vs_recall(precisions, recalls)  
plt.show()
```



```
In [28]: y_train_pred_90 = (y_scores > 70000)
```

```
In [29]: import warnings  
warnings.filterwarnings("ignore")  
precision_score(y_train_5, y_train_pred_90)
```

Out[29]: 0.0


```
In [30]: recall_score(y_train_5, y_train_pred_90)
```

Out[30]: 0.0

```
In [31]: from sklearn.metrics import roc_curve  
fpr, tpr, thresholds = roc_curve(y_train_5, y_scores)
```



```
In [32]: def plot_roc_curve(fpr, tpr, label=None):
plt.plot(fpr, tpr, linewidth=2, label=label)
plt.plot([0, 1], [0, 1], 'k--')
plt.axis([0, 1, 0, 1])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plot_roc_curve(fpr, tpr)
plt.show()
```




```
In [33]: from sklearn.metrics import roc_auc_score
roc_auc_score(y_train_5, y_scores)
```

Out[33]: 0.9606586694924489

```
In [34]: from sklearn.ensemble import RandomForestClassifier
forest_clf = RandomForestClassifier(random_state=42)
y_probas_forest = cross_val_predict(forest_clf, X_train, y_train_5, cv=3,
method="predict_proba")
```

```
In [35]: y_scores_forest = y_probas_forest[:, 1] # score = proba of positive class
fpr_forest, tpr_forest, thresholds_forest = roc_curve(y_train_5, y_scores_forest)
```

```
In [36]: plt.plot(fpr, tpr, "b:", label="SGD")
plot_roc_curve(fpr_forest, tpr_forest, "Random Forest")
plt.legend(loc="lower right")
plt.show()
```



```
In [37]: roc_auc_score(y_train_5, y_scores_forest)
```

```
Out[37]: 0.9984573419278522
```

```
In [38]: sgd_clf.fit(X_train, y_train) # y_train, not y_train_5
sgd_clf.predict([some_digit])
```

```
Out[38]: array([3], dtype=uint8)
```

```
In [39]: some_digit_scores = sgd_clf.decision_function([some_digit])
some_digit_scores
```

```
Out[39]: array([[ -38451.98135063, -27346.42156278, -30213.6866118 ,
    2354.83581704,  -522.53770922, -4292.47256279,
   -42491.8010252 , -13236.88812634, -3844.91582026,
   -2470.16305866]])
```

```
In [40]: np.argmax(some_digit_scores)
```

```
Out[40]: 3
```

```
In [41]: sgd_clf.classes_
```

```
Out[41]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)
```

```
In [42]: sgd_clf.classes_[5]
```

```
Out[42]: 5
```

```
In [43]: from sklearn.multiclass import OneVsOneClassifier
ovo_clf = OneVsOneClassifier(SGDClassifier(random_state=42))
ovo_clf.fit(X_train, y_train)
ovo_clf.predict([some_digit])
```

```
Out[43]: array([4], dtype=uint8)
```

```
In [44]: len(ovo_clf.estimators_)
```

```
Out[44]: 45
```

```
In [45]: forest_clf.fit(X_train, y_train)
forest_clf.predict([some_digit])
```

```
Out[45]: array([9], dtype=uint8)
```

```
In [46]: forest_clf.predict_proba([some_digit])
```

```
Out[46]: array([[0. , 0.02, 0. , 0. , 0.1 , 0.03, 0.02, 0.01, 0.01, 0.81]])
```

```
In [47]: cross_val_score(sgd_clf, X_train, y_train, cv=3, scoring="accuracy")
```

```
Out[47]: array([0.8833, 0.8761, 0.8477])
```


```
In [48]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train.astype(np.float64))
cross_val_score(sgd_clf, X_train_scaled, y_train, cv=3, scoring="accuracy")
```

```
Out[48]: array([0.90315, 0.9061 , 0.9034 ])
```

```
In [49]: y_train_pred = cross_val_predict(sgd_clf, X_train_scaled, y_train, cv=3)
conf_mx = confusion_matrix(y_train, y_train_pred)
conf_mx
```


```
Out[49]: array([[5604,  0, 16,  6,  9, 50, 35,  6, 196,  1],
 [  1, 6425, 44, 20,  4, 49,  5,  8, 173, 13],
 [ 28,  28, 5284, 87, 72, 25, 69, 38, 317, 10],
 [ 28,  22, 111, 5260,  2, 223, 25, 42, 350, 68],
 [  7, 12, 44, 11, 5259, 11, 43, 20, 263, 172],
 [ 28, 17, 27, 155, 55, 4551, 77, 15, 428, 68],
 [ 27, 18, 55,  3, 43, 95, 5556,  5, 116,  0],
 [ 18, 14, 51, 22, 50, 17,  3, 5717, 144, 229],
 [ 17, 66, 47, 88,  2, 131, 36,  7, 5406, 51],
 [ 23, 23, 28, 63, 126, 41,  1, 171, 282, 5191]],
 dtype=int64)
```

```
In [50]: plt.matshow(conf_mx, cmap=plt.cm.gray)
plt.show()
#可见分类效果非常的好
```




```
In [51]: row_sums = conf_mx.sum(axis=1, keepdims=True)
norm_conf_mx = conf_mx / row_sums
```

```
In [52]: np.fill_diagonal(norm_conf_mx, 0)
plt.matshow(norm_conf_mx, cmap=plt.cm.gray)
plt.show()
#分析误差，其他数字被误分类成8的可能性最大
```



```
In [53]: cl_a, cl_b = 3, 5
X_aa = X_train[(y_train == cl_a) & (y_train_pred == cl_a)]
X_ab = X_train[(y_train == cl_a) & (y_train_pred == cl_b)]
X_ba = X_train[(y_train == cl_b) & (y_train_pred == cl_a)]
X_bb = X_train[(y_train == cl_b) & (y_train_pred == cl_b)]
plt.figure(figsize=(8,8))
plt.subplot(221); plot_digits(X_aa[:25], images_per_row=5)
plt.subplot(222); plot_digits(X_ab[:25], images_per_row=5)
plt.subplot(223); plot_digits(X_ba[:25], images_per_row=5)
plt.subplot(224); plot_digits(X_bb[:25], images_per_row=5)
plt.show()
```



```
In [54]: from sklearn.neighbors import KNeighborsClassifier
y_train_large = (y_train >= 7)
y_train_odd = (y_train % 2 == 1)
y_multilabel = np.c_[y_train_large, y_train_odd]
knn_clf = KNeighborsClassifier()
knn_clf.fit(X_train, y_multilabel)
```

```
Out[54]: ▼ KNeighborsClassifier
KNeighborsClassifier()
```

```
In [55]: knn_clf.predict([some_digit])
```



```
Out[55]: array([[ True,  True]])
```

```
In [56]: y_train_knn_pred = cross_val_predict(knn_clf, X_train, y_train, cv=3)
f1_score(y_train, y_train_knn_pred, average="macro")
```


```
Out[56]: 0.96801255083268
```

```
In [57]: noise = np.random.randint(0, 100, (len(X_train), 784))
X_train_mod = X_train + noise
noise = np.random.randint(0, 100, (len(X_test), 784))
X_test_mod = X_test + noise
y_train_mod = X_train
y_test_mod = X_test
```

```
In [58]: some_index = 5500
plt.subplot(121); plot_digit(X_test_mod[some_index])
plt.subplot(122); plot_digit(y_test_mod[some_index])
plt.show()
```



```
In [59]: knn_clf.fit(X_train_mod, y_train_mod)
clean_digit = knn_clf.predict([X_test_mod[some_index]])
plot_digit(clean_digit)
```



```
In [ ]:
```